

SciNumTools DIPH5

Environment HDF5 Format Specification

Version 2.5

Contents

1	Scope	1
2	Persisted environment boundary	1
3	File identification and versioning	2
3.1	File naming	2
4	Paths and objects	2
4.1	Collections	2
5	Value datasets	3
5.1	Value type table	3
6	Node settings and provenance	3
6.1	Source manifest	4
6.2	Trace manifest	4
7	Examples	4
7.1	Scalar values and hierarchy	4
7.2	Arrays and units	5
7.3	Value groups	5
7.4	Custom-unit manifest	5
7.5	Ordered collections	6
8	Compatibility	6

1 Scope

DIPH5 is the HDF5 interchange format used by SciNumTools to persist an *evaluated DIPL environment*. It stores evaluated value nodes together with their DIPL paths, collection structure, units, settings, and provenance.

The persisted object is therefore a data environment suitable for exchange with scientific applications. It is not a serialization of the complete parser or evaluation runtime.

This document specifies the on-disk representation. It does not define the DIPL language itself; the DIPL specification defines the language concepts represented here.

2 Persisted environment boundary

The file contains the evaluated value-node tree and the node-level information needed to interpret and trace those values. In particular, node settings, units, schemas, source information, and provenance are persisted as dataset attributes.

Implementation-specific runtime registries are not serialized as executable objects. This includes global function definitions, branching state, and the complete unit and schema registries. DIPH5 version 2 additionally stores a source manifest containing source identities, paths, parent relationships, and content fingerprints. Version 2.1 also stores durable identifiers for registered units, schemas, and functions. It does not embed complete

source text, parsed source nodes, or executable functions. Version 2.2 adds `value_group` objects. Version 2.3 persists custom PUEL unit registrations required to interpret evaluated quantities. Version 2.4 adds descriptive schema metadata and source locations to schema trace entries. Version 2.5 records which evaluated values were overridden and where their replacement values originated. A loaded DIPH5 file must therefore be treated as an evaluated environment, not as a source from which the original DIPL program can be reconstructed exactly.

3 File identification and versioning

The HDF5 root object MUST contain the following scalar attributes:

Attribute	Type	Required value
<code>_DIPL_Format</code>	UTF-8 string	SciNumTools3 Environment
<code>_DIPL_Schema_Version</code>	unsigned integer	2
<code>_DIPL_Schema_Version_Minor</code>	unsigned integer	5

Readers MUST reject files with a different format identifier or unsupported schema version. Version 2.5 readers support version 1 files and versions 2.0 through 2.5. Future schema revisions MUST preserve the meaning of existing attributes or increment the major or minor schema version.

3.1 File naming

Files SHOULD use the `.diph5` filename extension to make their purpose clear. The extension is only a naming convention and is not normative. Readers MUST identify a file from the root `_DIPL_Format` and `_DIPL_Schema_Version` attributes, regardless of its filename.

4 Paths and objects

Each component of a fully qualified DIPL path is represented by an HDF5 path component. A value node without child nodes is an HDF5 dataset at its resolved path. Intermediate groups represent DIPL hierarchy and carry `_DIPL_Kind=group` and `_DIPL_Path` attributes.

The `_DIPL_Path` attribute stores the corresponding fully qualified DIPL path and is the authoritative mapping back to DIPL semantics.

4.1 Collections

Collection containers and items are represented by groups. Their meaning is identified by `_DIPL_Kind`:

<code>_DIPL_Kind</code>	Meaning	Additional attributes
<code>map</code>	keyed collection	—
<code>map_item</code>	item of a keyed collection	<code>_DIPL_Key</code>
<code>list</code>	ordered collection	—
<code>list_item</code>	item of an ordered collection	<code>_DIPL_Index</code>

Collection item groups may contain nested groups and value datasets. They are containers and cannot themselves represent DIPL value nodes.

The complete object-kind table is:

HDF5 object	<code>_DIPL_Kind</code>	<code>_DIPL_Path</code>	Other required attributes
hierarchy group	<code>group</code>	group path	—
keyed collection	<code>map</code>	collection path	—
keyed item	<code>map_item</code>	item path	<code>_DIPL_Key</code>
ordered collection	<code>list</code>	collection path	—
ordered item	<code>list_item</code>	item path	<code>_DIPL_Index</code>
value with children	<code>value_group</code>	value path	<code>_DIPL_Value</code> dataset

HDF5 object	<code>_DIPL_Kind</code>	<code>_DIPL_Path</code>	Other required attributes
value dataset	<code>value</code>	value path	<code>_DIPL_Node_Type</code> , <code>_DIPL_Value_Type</code>

An implementation **MUST** use the kind attribute to distinguish a collection from an ordinary hierarchy group. HDF5 object names alone are not sufficient, because a numbered or keyed collection item is deliberately flattened into a normal group path.

5 Value datasets

Every value dataset has `_DIPL_Kind=value`, `_DIPL_Path`, `_DIPL_Node_Type`, and `_DIPL_Value_Type` attributes. `_DIPL_Array` identifies array values; scalar values use a scalar dataspace and arrays use a simple dataspace whose dimensions match the DIPL value shape.

Values are written using native HDF5 numeric types. Strings use fixed-width UTF-8 HDF5 string datatypes. A null DIPL value is represented by a dataset with an HDF5 null dataspace.

5.1 Value type table

The `_DIPL_Value_Type` value identifies the logical DIPL type. The following table lists the concrete types currently written by DIP5 and their HDF5 storage types:

<code>_DIPL_Value_Type</code>	Logical value	HDF5 datatype	Scalar/array
<code>bool</code>	Boolean	native unsigned 8-bit integer	both
<code>char</code>	Character	native signed 8-bit integer	both
<code>int8</code> , <code>int16</code> , <code>int32</code> , <code>int64</code>	Signed integer	matching native signed integer	both
<code>uint8</code> , <code>uint16</code> , <code>uint32</code> , <code>uint64</code>	Unsigned integer	matching native unsigned integer	both
<code>float32</code>	Single precision real	native 32-bit float	both
<code>float64</code>	Double precision real	native 64-bit float	both
<code>float128</code>	Extended precision real	native long double	both
<code>str</code>	UTF-8 string	fixed-width UTF-8 string	both

For a scalar, the dataset has an HDF5 scalar dataspace. For an array, the dataset has a simple dataspace and `_DIPL_Array` is non-zero. The dimensions are stored in the dataspace rather than in a separate attribute.

6 Node settings and provenance

When present, node settings are stored as attributes on the value dataset:

`_DIPL_Constant`, `_DIPL_Condition`, `_DIPL_Format_Spec`, `_DIPL_Tags`, `_DIPL_Options`, `_DIPL_Option_Units`, `_DIPL_Schemas`, and `_DIPL_Value_Origin`.

Version 2.5 adds `_DIPL_Override` (unsigned 8-bit boolean) to each value dataset. If true, `_DIPL_Override_Source`, `_DIPL_Override_Line`, and `_DIPL_Override_Code` identify the modification that supplied the effective value. The ordinary source attributes continue to identify the declaration.

Source and provenance information uses `_DIPL_Source`, `_DIPL_Source_Line`, `_DIPL_Source_Code`, and the metadata attributes `description`, `authors`, `title`, `journal`, `year`, `volume`, `issue`, `pages`, `doi`, `url`, `version`, `created`, `modified`, `license`, `rationale`, `recommended_range`, `performance_impact`, `scientific_impact`, `deprecated`, `replacement`, `since`, `category`, and `visibility`. The plural metadata attributes `native`, `requires`, `conflicts`, `implies`, `see`, and `example` are UTF-8 string-array attributes. Physical units are stored in the `units` attribute using their canonical PUEL representation.

6.1 Source manifest

Version 2 files contain a root `/_DIPL_Sources` group. Each child is a numbered source-manifest entry so DIPL source names never need to be encoded as HDF5 object names. Each entry has the following attributes:

Attribute	Type	Meaning
<code>_DIPL_Source_Name</code>	UTF-8 string	Named DIPL source identifier.
<code>_DIPL_Source_Path</code>	UTF-8 string	Path recorded when the source was parsed; it may be empty for inline input.
<code>_DIPL_Source_Parent</code>	UTF-8 string	Parent source identifier, when one exists.
<code>_DIPL_Source_Parent_Line</code>	unsigned integer	Parent source line associated with the registration.
<code>_DIPL_Source_Hash_Algorithm</code>	UTF-8 string	Digest algorithm, currently SHA-256 .
<code>_DIPL_Source_Hash</code>	UTF-8 string	Lowercase hexadecimal digest of the exact UTF-8 bytes parsed by DIP.

The manifest is an integrity and provenance aid, not a source archive. A consumer can hash an available source file and compare the result with the stored digest, but loading DIPH5 does not recreate executable source definitions or enable source-qualified lookups. `_DIPL_Sources` is reserved at the DIPH5 root; a DIPL environment with that top-level path cannot be saved.

6.2 Trace manifest

Version 2.1 files contain a root `/_DIPL_Trace` group. Each child is a numbered trace-manifest entry. It records the stable internal identifier assigned during parsing for every registered unit, schema, and function:

Attribute	Type	Meaning
<code>_DIPL_Trace_Id</code>	UTF-8 string	Internal trace identifier, for example <code>DIP0_UNIT0</code> .
<code>_DIPL_Trace_Name</code>	UTF-8 string	Registered public name.
<code>_DIPL_Trace_Kind</code>	UTF-8 string	unit , schema , function_value , or function_nodes .

The trace manifest preserves diagnostic and documentation identity only. It does not recreate a custom unit definition, schema nodes, or an executable function on load. `_DIPL_Trace` is reserved at the DIPH5 root; a DIPL environment with that top-level path cannot be saved.

In version 2.4, schema trace entries additionally carry `_DIPL_Schema_Source` (the source-manifest name), `_DIPL_Schema_Source_Line` (unsigned declaration or first-body line), and the same optional metadata attributes used by value nodes, including **description**, **authors**, and **doi**. These fields describe the schema and its origin; they do not provide a reusable definition. Earlier files retain schema IDs and names but have no schema-level metadata or source location. Schema-derived value datasets also carry `_DIPL_Node_Schema_Id`, which refers to the schema trace ID that supplied the node. Group and collection schema associations remain in `_DIPL_Schemas` attributes on their HDF5 groups.

7 Examples

7.1 Scalar values and hierarchy

This DIPL fragment:

```
simulation
  steps int = 100
  enabled bool = true
```

is represented as:

```

/simulation          group  (_DIPL_Kind=group)
/simulation/steps    dataset (_DIPL_Kind=value,
                        _DIPL_Path=simulation.steps,
                        _DIPL_Value_Type=int32)
/simulation/enabled  dataset (_DIPL_Kind=value,
                        _DIPL_Path=simulation.enabled,
                        _DIPL_Value_Type=bool)

```

`/simulation/steps` has a scalar dataspace and a native signed 32-bit HDF5 datatype. `/simulation/enabled` has a scalar dataspace and a native unsigned 8-bit HDF5 datatype.

7.2 Arrays and units

This DIPL value:

```

boundary[inlet]
  velocity float[3] = [1.0, 0.0, 0.0] m/s

```

is represented as:

```

/boundary          group  (_DIPL_Kind=map,
                        _DIPL_Path=boundary)
/boundary/inlet    group  (_DIPL_Kind=map_item,
                        _DIPL_Path=boundary[inlet],
                        _DIPL_Key=inlet)
/boundary/inlet/velocity dataset (_DIPL_Kind=value,
                        _DIPL_Path=boundary[inlet].velocity,
                        _DIPL_Value_Type=float64,
                        _DIPL_Array=1,
                        units=m*s-1)

```

The velocity dataset has rank one and dimension 3.

7.3 Value groups

When a DIPL value node has descendants, its resolved HDF5 path is a group rather than a dataset. The group has `_DIPL_Kind=value_group` and `_DIPL_Path` set to the value node's fully qualified path. Its own value is stored in a reserved child dataset named `_DIPL_Value`; that dataset uses the normal `value` dataset attributes, including the real DIPL path rather than a synthetic sentinel path. All other children of the group are regular DIPL descendants.

```

feature bool = true
  setting int = 2

```

is represented as:

```

/feature          group  (_DIPL_Kind=value_group,
                        _DIPL_Path=feature)
/feature/_DIPL_Value dataset (_DIPL_Kind=value,
                        _DIPL_Path=feature,
                        _DIPL_Value_Type=bool)
/feature/setting  dataset (_DIPL_Kind=value,
                        _DIPL_Path=feature.setting,
                        _DIPL_Value_Type=int32)

```

`_DIPL_Value` is format-owned and is not a DIPL child node. A direct DIPL child with that name beneath a value group is invalid and writers MUST reject it. The payload is a value node and therefore MUST NOT carry applied schemas; schemas are represented only by their expanded group or collection-item structure.

7.4 Custom-unit manifest

DIP5 2.3 stores custom PUEL registrations in `/_DIPL_Units`. Each numbered entry carries `_DIPL_Unit_Name`, `_DIPL_Unit_Definition`, `_DIPL_Unit_Order`, and `_DIPL_Trace_Id`. Readers MUST register entries in ascending order before reading value datasets, so definitions may depend on earlier custom units. `_DIPL_Units` is format-owned and cannot be a top-level DIPL path.

7.5 Ordered collections

This DIPL fragment:

```
samples[]  
  time float = 0 s  
samples[]  
  time float = 1 s
```

is represented by a list group containing list_item groups:

```
/samples          group  (_DIPL_Kind=list,  
                        _DIPL_Path=samples)  
/samples/0        group  (_DIPL_Kind=list_item,  
                        _DIPL_Path=samples[0],  
                        _DIPL_Index=0)  
/samples/0/time   dataset (_DIPL_Path=samples[0].time)  
/samples/1        group  (_DIPL_Kind=list_item,  
                        _DIPL_Path=samples[1],  
                        _DIPL_Index=1)  
/samples/1/time   dataset (_DIPL_Path=samples[1].time)
```

The `_DIPL_Index` attribute, rather than HDF5 group iteration order, defines the order of list items.

8 Compatibility

Implementations SHOULD preserve unknown attributes when editing a DIP5 file and MUST ignore attributes they do not understand. They MUST reject a file whose format identifier or schema version is unsupported rather than silently interpreting it as a different format.